

# aerofit-project-1

August 14, 2023

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
```

```
[2]: # Importing the dataset and basic df analysis
df = pd.read_csv(r'aerofit_treadmill.csv')
```

```
[3]: # first few rows of the dataset
df.head()
```

```
[3]:   Product  Age  Gender  Education  MaritalStatus  Usage  Fitness  Income  Miles
0   KP281   18   Male       14         Single        3        4   29562   112
1   KP281   19   Male       15         Single        2        3   31836    75
2   KP281   19  Female       14   Partnered        4        3   30699    66
3   KP281   19   Male       12         Single        3        3   32973    85
4   KP281   20   Male       13   Partnered        4        2   35247    47
```

```
[4]: # Shape of dataset
df.shape
```

```
[4]: (180, 9)
```

```
[5]: # basic information about the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 180 entries, 0 to 179
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Product         180 non-null   object
1   Age             180 non-null   int64
2   Gender          180 non-null   object
3   Education       180 non-null   int64
4   MaritalStatus   180 non-null   object
5   Usage           180 non-null   int64
6   Fitness         180 non-null   int64
```

```

7   Income      180 non-null   int64
8   Miles       180 non-null   int64
dtypes: int64(6), object(3)
memory usage: 12.8+ KB

```

```

[6]: # statistics of the numerical columns
df.describe(include="all").T.round(1)

```

```

[6]:
count unique      top freq      mean      std \
Product      180      3      KP281  80      NaN      NaN
Age          180.0    NaN      NaN  NaN      28.788889    6.943498
Gender       180      2      Male  104      NaN      NaN
Education    180.0    NaN      NaN  NaN      15.572222    1.617055
MaritalStatus 180      2  Partnered  107      NaN      NaN
Usage        180.0    NaN      NaN  NaN      3.455556    1.084797
Fitness      180.0    NaN      NaN  NaN      3.311111    0.958869
Income       180.0    NaN      NaN  NaN      53719.577778  16506.684226
Miles        180.0    NaN      NaN  NaN      103.194444    51.863605

      min      25%      50%      75%      max
Product      NaN      NaN      NaN      NaN      NaN
Age          18.0     24.0     26.0     33.0     50.0
Gender       NaN      NaN      NaN      NaN      NaN
Education     12.0     14.0     16.0     16.0     21.0
MaritalStatus  NaN      NaN      NaN      NaN      NaN
Usage         2.0      3.0      3.0      4.0      7.0
Fitness       1.0      3.0      3.0      4.0      5.0
Income      29562.0  44058.75  50596.5  58668.0  104581.0
Miles        21.0     66.0     94.0    114.75    360.0

```

```

[7]: #
df.describe()

```

```

[7]:
count      Age      Education      Usage      Fitness      Income \
count  180.000000  180.000000  180.000000  180.000000    180.000000
mean    28.788889   15.572222   3.455556   3.311111   53719.577778
std      6.943498   1.617055   1.084797   0.958869   16506.684226
min     18.000000   12.000000   2.000000   1.000000   29562.000000
25%     24.000000   14.000000   3.000000   3.000000   44058.750000
50%     26.000000   16.000000   3.000000   3.000000   50596.500000
75%     33.000000   16.000000   4.000000   4.000000   58668.000000
max     50.000000   21.000000   7.000000   5.000000  104581.000000

      Miles
count  180.000000
mean   103.194444
std     51.863605

```

```
min      21.000000
25%      66.000000
50%      94.000000
75%     114.750000
max     360.000000
```

```
[8]: df.Education.value_counts()
```

```
[8]: 16      85
      14      55
      18      23
      15       5
      13       5
      12       3
      21       3
      20       1
      Name: Education, dtype: int64
```

```
[9]: df.Age.value_counts()
```

```
[9]: 25      25
      23      18
      24      12
      26      12
      28       9
      35       8
      33       8
      30       7
      38       7
      21       7
      22       7
      27       7
      31       6
      34       6
      29       6
      20       5
      40       5
      32       4
      19       4
      48       2
      37       2
      45       2
      47       2
      46       1
      50       1
      18       1
      44       1
```

```

43      1
41      1
39      1
36      1
42      1
Name: Age, dtype: int64

```

```
[10]: df.Product.value_counts()
```

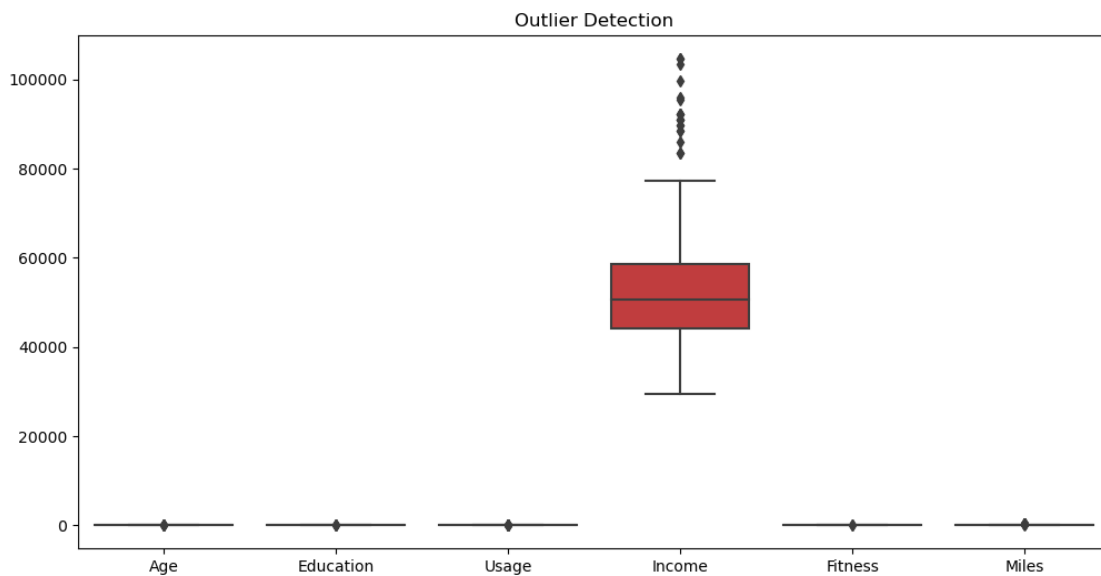
```

[10]: KP281      80
      KP481      60
      KP781      40
      Name: Product, dtype: int64

```


```
[11]: categorical_cols = ['Product', 'Gender', 'MaritalStatus']
      df[categorical_cols] = df[categorical_cols].astype('category')
```

```
[12]: # Detecting Outliers
      # identify and handle outliers using box plots
      numeric_cols = ['Age', 'Education', 'Usage', 'Income', 'Fitness', 'Miles']
      plt.figure(figsize=(12, 6))
      sns.boxplot(data=df[numeric_cols])
      plt.title("Outlier Detection")
      plt.show()
```

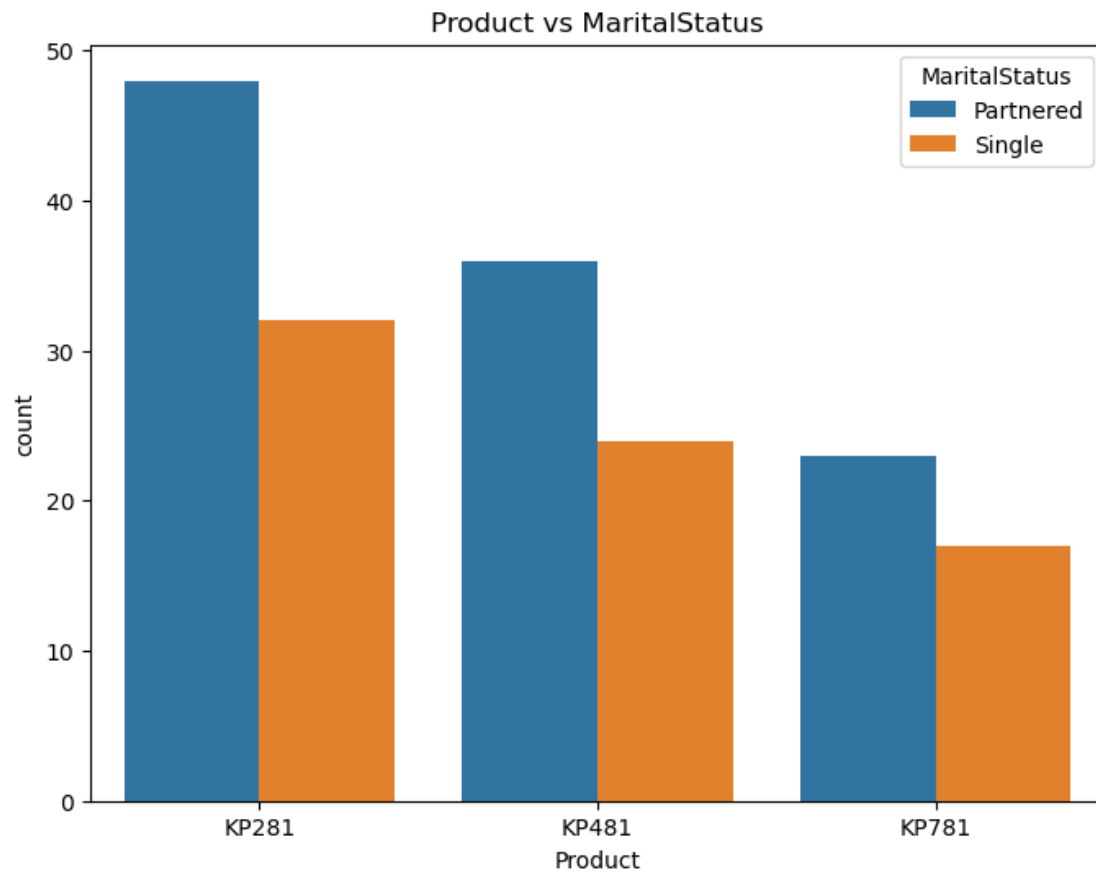


```
[13]: # Visual Analysis - Univariate & Bivariate
      # Histogram of Age
      plt.figure(figsize=(8, 6))
```


```
sns.histplot(data=df, x='Age', bins=20, kde=True)
plt.title("Distribution of Age")
plt.show()
```



```
[14]: # Countplot for Product vs MaritalStatus
plt.figure(figsize=(8, 6))
sns.countplot(x="Product", hue='MaritalStatus', data=df)
plt.title("Product vs MaritalStatus")
plt.show()
```



```
[15]: # Boxplot of Product vs Age
plt.figure(figsize=(8, 6))
sns.boxplot(data=df, x='Product', y='Age')
plt.title("Product vs Age")
plt.show()
```



```
[16]: # Representing marginal probability
# Marginal probabilities of purchasing each product
marginal_prob = pd.crosstab(index=df['Product'], columns='count',
                             ↪normalize=True)
print("Marginal Probability:")
print(marginal_prob)
```

Marginal Probability:

col_0	count
-------	-------

Product
---------

KP281	0.444444
-------	----------

KP481	0.333333
-------	----------

KP781	0.222222
-------	----------

```
[17]: # Correlation Analysis
# correlation matrix
corr = df.corr()
```

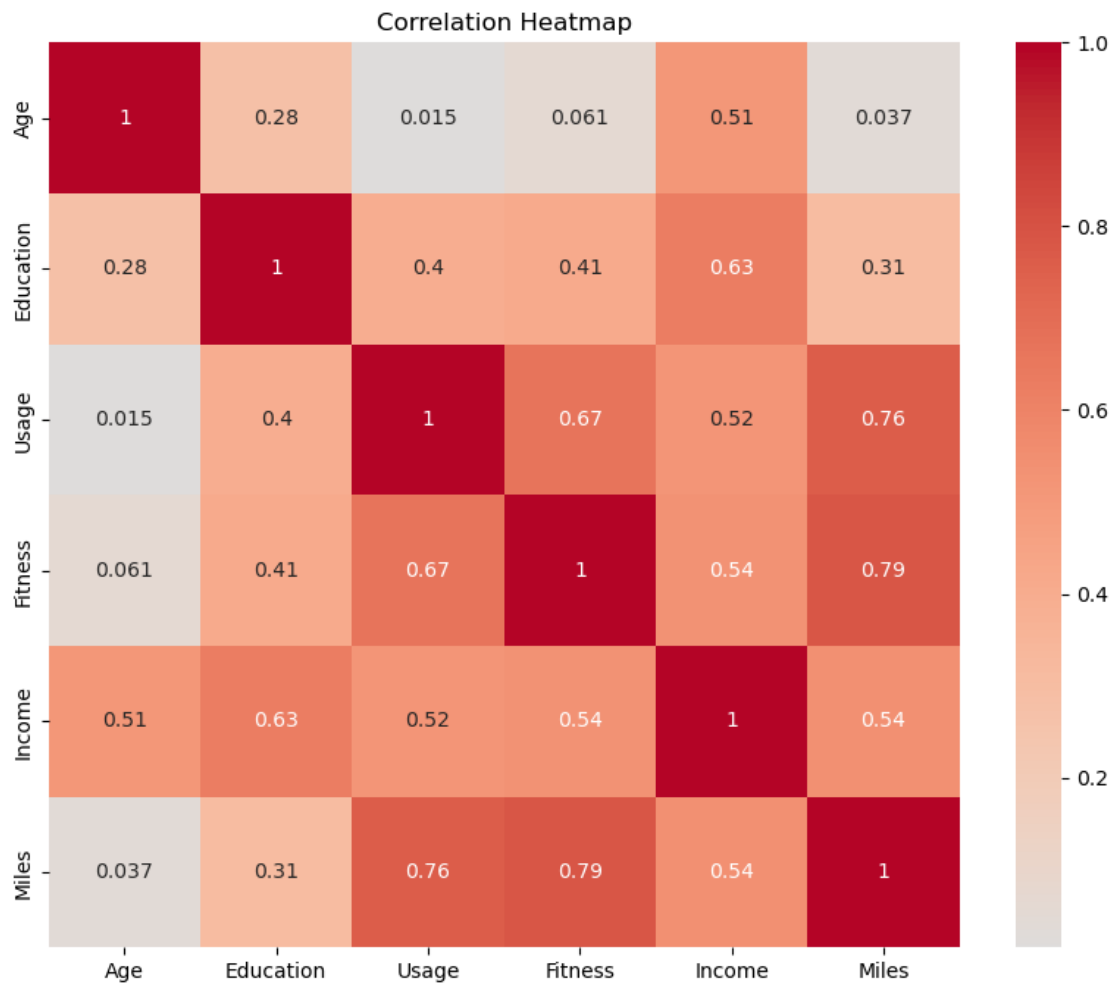
C:\Users\rahul\AppData\Local\Temp\ipykernel\_10748\2035915871.py:3:

FutureWarning: The default value of numeric\_only in DataFrame.corr is

deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
corr = df.corr()
```

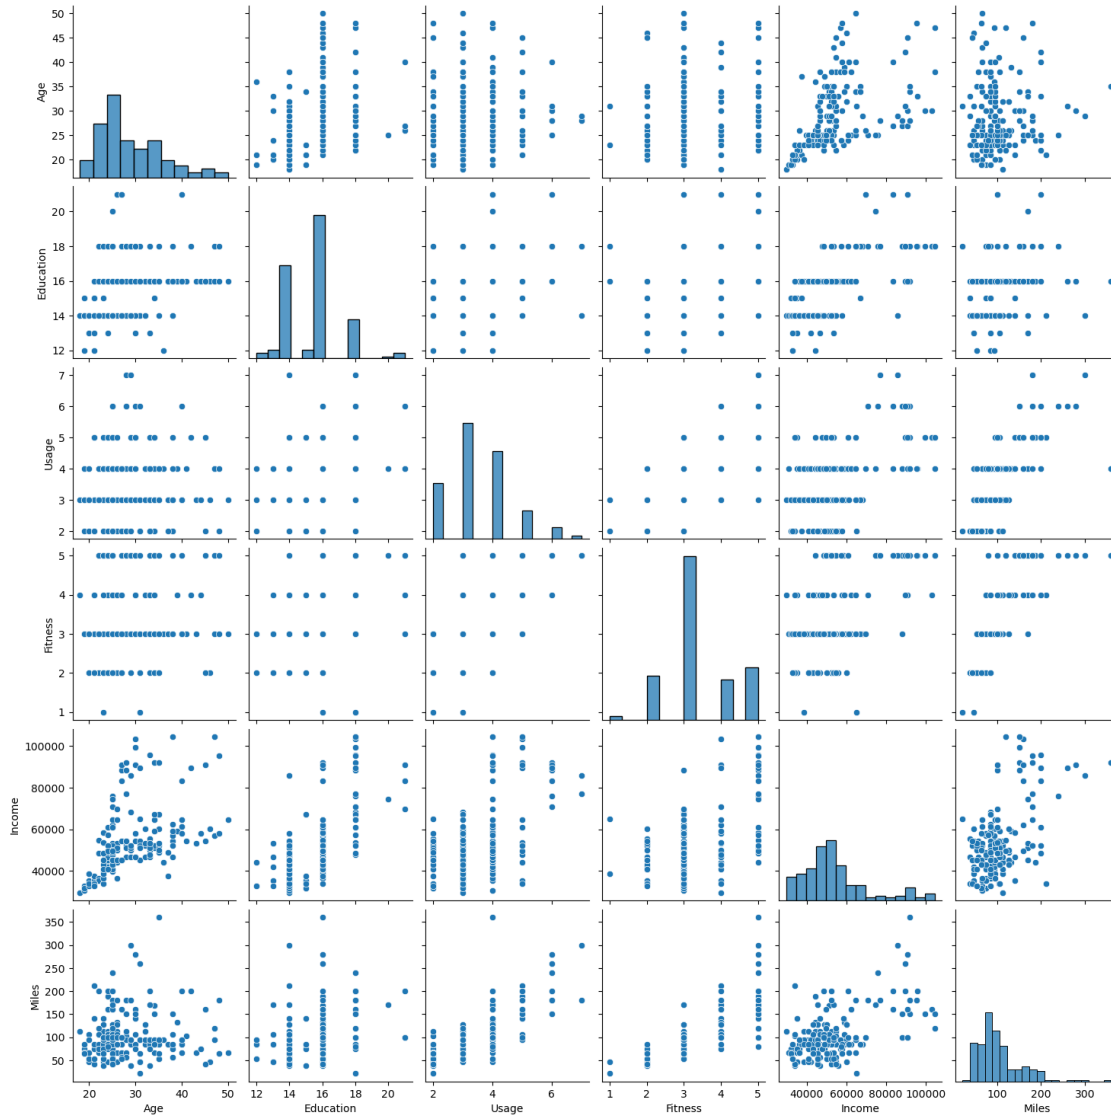
```
[18]: # Heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, cmap='coolwarm', center=0)
plt.title("Correlation Heatmap")
plt.show()
```



```
[19]: # Pairplot
sns.pairplot(df)
```

```
[19]: <seaborn.axisgrid.PairGrid at 0x18ec25a3d00>
```





```
[20]: # Calculating conditional probability
df.std()
```

```
C:\Users\rahul\AppData\Local\Temp\ipykernel_10748\3187971767.py:2:
FutureWarning: The default value of numeric_only in DataFrame.std is deprecated.
In a future version, it will default to False. In addition, specifying
'numeric_only=None' is deprecated. Select only valid columns or specify the
value of numeric_only to silence this warning.
df.std()
```

```
[20]: Age                6.943498
      Education          1.617055
      Usage              1.084797
```

```
Fitness          0.958869
Income          16506.684226
Miles           51.863605
dtype: float64
```

```
[21]: # Median
df.median()
```

```
C:\Users\rahul\AppData\Local\Temp\ipykernel_10748\3213049308.py:2:
FutureWarning: The default value of numeric_only in DataFrame.median is
deprecated. In a future version, it will default to False. In addition,
specifying 'numeric_only=None' is deprecated. Select only valid columns or
specify the value of numeric_only to silence this warning.
    df.median()
```

```
[21]: Age          26.0
      Education    16.0
      Usage        3.0
      Fitness      3.0
      Income      50596.5
      Miles       94.0
      dtype: float64
```

```
[22]: # Mean
df.mean()
```

```
C:\Users\rahul\AppData\Local\Temp\ipykernel_10748\2486691740.py:2:
FutureWarning: The default value of numeric_only in DataFrame.mean is
deprecated. In a future version, it will default to False. In addition,
specifying 'numeric_only=None' is deprecated. Select only valid columns or
specify the value of numeric_only to silence this warning.
    df.mean()
```

```
[22]: Age          28.788889
      Education    15.572222
      Usage        3.455556
      Fitness      3.311111
      Income      53719.577778
      Miles       103.194444
      dtype: float64
```

```
[23]: age_diff = df['Age'].mean() - df['Age'].median()
      print(f'Difference between mean and median of Age: {age_diff}')
```

```
Difference between mean and median of Age: 2.7888888888888888
```

```
[24]: # Gender for each Product
conditional_prob_gender = pd.crosstab(index=df['Product'],
    ↪ columns=df['Gender'], normalize='index')
print("Conditional Probability (Gender):")
print(conditional_prob_gender)
```

```
Conditional Probability (Gender):
Gender      Female      Male
Product
KP281      0.500000  0.500000
KP481      0.483333  0.516667
KP781      0.175000  0.825000
```

```
[25]: # Customer Profiling - Categorization of users
X = df[['Age', 'Education', 'Usage', 'Income', 'Fitness', 'Miles']]
```

```
[26]: # KMeans clustering
kmeans = KMeans(n_clusters=3, random_state=0)
df['Cluster'] = kmeans.fit_predict(X)
```

```
C:\ProgramData\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\ProgramData\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
    warnings.warn(
```

```
[27]: # Generating recommendations and actionable insights
# Customer profiling insights
for cluster in range(3):
    cluster_df = df[df['Cluster'] == cluster]
    print(f"Cluster {cluster} has {len(cluster_df)} customers.")
    print(f"Average age: {cluster_df['Age'].mean():.2f}")
    print(f"Average income: ${cluster_df['Income'].mean():.2f}")
    print(f"Average fitness rating: {cluster_df['Fitness'].mean():.2f}")
    print()
```

```
Cluster 0 has 22 customers.
Average age: 33.05
Average income: $90,417.00
Average fitness rating: 4.77
```

```
Cluster 1 has 85 customers.
Average age: 31.41
Average income: $55,564.86
```

Average fitness rating: 3.25

Cluster 2 has 73 customers.

Average age: 24.45

Average income: \$40,511.47

Average fitness rating: 2.95

```
[28]: # Recommendations
recommendations = [
    "Cluster 0 consists of customers with relatively lower age, income, and
    ↪fitness rating.Consider targeting them with entry-level products like KP281.
    ↪",
    "Cluster 1 comprises customers with higher income and fitness levels. They
    ↪might be interested in advanced treadmills like KP781.",
    "Cluster 2 represents a diverse group. It could be further analyzed to
    ↪identify specific customer segments and tailor marketing strategies
    ↪accordingly."
]

for i, recommendation in enumerate(recommendations):
    print(f"Recommendation {i + 1}: {recommendation}")
```

Recommendation 1: Cluster 0 consists of customers with relatively lower age, income, and fitness rating.Consider targeting them with entry-level products like KP281.

Recommendation 2: Cluster 1 comprises customers with higher income and fitness levels. They might be interested in advanced treadmills like KP781.

Recommendation 3: Cluster 2 represents a diverse group. It could be further analyzed to identify specific customer segments and tailor marketing strategies accordingly.

```
[ ]:
```